

Kevin Yang

New York, NY | +1(347)278-9339 | ky2738@nyu.edu | [LinkedIn](#) | [Portfolio](#) | [GitHub](#)

EDUCATION

New York University

Master of Science in Computer Engineering

- Concentration: AI& Data Modeling, GPA: 3.8/4.0

New York, NY

Sep. 2025 – May 2027

The Chinese University of Hong Kong

Bachelor of Engineering in Computer Engineering

- Top 20% of program

Hong Kong

Sep. 2021 – Jun. 2025

PROFESSIONAL EXPERIENCE

Yanfeng Automotive Interiors

Software Engineer Intern

Novi, MI

May 2026 – Aug. 2026

- Architected and shipped a multi-agent decision platform used daily by operations, finance, and planning teams (**Python** / **FastAPI** microservices, **OpenAI SDK**, **LangGraph**), cutting data consolidation from **2 weeks to 10 min per cycle**
- Built the **Data Agent** on **PostgreSQL** with a typed semantic layer and hybrid **BM25 + vector** retrieval (RRF fusion, cross-encoder reranking), raising correct-evidence retrieval from **58% to 90% (Recall@10)**
- Designed a fault-tolerant **Agent Harness** on **PostgreSQL**, **Redis**, and **Temporal** (durable, idempotent, resumable workflows; four-tier memory; token-aware context compression), reducing P95 context-assembly latency from **1.2s to 430ms**
- Added end-to-end observability (**Langfuse** tracing) and a CI/CD-style release gate that improved testability: Skill candidates from user corrections and traced failures must pass golden-set regression, sandbox replay, canary release, and auto-rollback, cutting regressions on solved tasks from **21% to 5%**

Capgemini

Software Engineer Intern

Shanghai, China

Aug. 2024 – Sep. 2024

- Built a high-concurrency flash-sale service in **Java** / **Spring Boot** with **Redis Lua** atomic reservations and **RocketMQ** buffering: **1,000+ QPS at P99 < 300 ms**, zero oversells or duplicates across **12 fault-injection scenarios**
- Designed a resilient **Redis** / **RocketMQ** / **MySQL** consistency layer (transactional outbox, idempotent consumers, cross-domain reconciliation), recovering from **95% of database failures** and cutting recovery time from **30 to 3 minutes**
- Root-caused a **Tomcat bottleneck** and a **MyBatis-Plus concurrency bug** through load testing, log analysis, and an **AI-assisted inspection tool** (CPU spikes, memory leaks, OOMs); added a multi-tier **Redis** cache with hot-key protection, cutting peak MySQL reads by **75%** and P99 latency by **60%**
- Designed an annotation-driven (**@TaskNode**) orchestration engine and a declarative **state-machine** component with failure / retry policies and audit logging, using an MVCC-inspired **copy-merge context model** to eliminate data races across parallel nodes; cut multi-source aggregation latency by **50%**, with new workflows onboarding at zero core-engine changes

NIO

Machine Learning Engineer Intern

Shanghai, China

May 2024 – Jul. 2024

- Formulated defect detection as a preference-optimization problem: built a pairwise preference dataset with senior inspectors and trained **LoRA (r=16) SFT + DPO** in **PyTorch** (Bradley-Terry objective), raising detection accuracy from **76% to 91%**
- Rebuilt the evaluation pipeline around a refreshed golden set and an **OpenCV** / **YOLO** pre-screening model, tuning thresholds on precision/recall trade-offs to raise annotation agreement from **71% to 94%** and cut labeling effort by **55%**

SKILLS

- Languages:** Python | Java | SQL | TypeScript | JavaScript | Bash
- AI & LLM:** LLM application development | RAG | Multi-agent systems | AI agents | OpenAI API | LangChain | LangGraph | Hugging Face | Fine-tuning (LoRA, SFT, DPO) | Prompt engineering | Embeddings & vector search | Text-to-SQL | LLM evaluation
- Machine Learning:** PyTorch | scikit-learn | NumPy | Pandas | Computer vision | YOLO | OpenCV | TensorRT | Model training & deployment
- Backend:** FastAPI | Spring Boot | REST APIs | Microservices | Distributed systems | High-concurrency systems | Asynchronous programming
- Data & Messaging:** PostgreSQL | MySQL | Redis | Elasticsearch | pgvector | RocketMQ | Message queues | Object storage (MinIO, S3) | Data pipelines
- Cloud & DevOps:** Docker | Kubernetes | Azure | Linux | Git | GitHub Actions | CI/CD | Monitoring & logging