

Kevin Yang

New York, NY | +1(347)278-9339 | ky2738@nyu.edu | [LinkedIn](#) | [Portfolio](#) | [GitHub](#)

EDUCATION

New York University

New York, NY

Master of science in Computer Engineering

Sep. 2025 – May 2027

- Concentration: AI& Data Modeling, GPA: 3.8/4.0

The Chinese University of Hong Kong

Hong Kong

Bachelor of Engineering in Computer Engineering

Sep. 2021 – Jun. 2025

- Top 20% of program

PROFESSIONAL EXPERIENCE

Yanfeng Automotive Interiors

Novi, MI

Software Engineer Intern

May 2026 – Aug. 2026

- Architected and built a multi-agent decision platform for corporate strategy using **Python, FastAPI, OpenAI SDK, LangGraph**, cutting source data consolidation, reconciliation, and governance from **2 weeks to 10 min per cycle** at **95% figure-level accuracy**
- Built the **Data Agent** on **PostgreSQL** with a semantic layer and adopted vector + keyword retrieval to raise correct-evidence retrieval from **58% to 90% (Recall@10)**
- Created a new **Agent Harness** layer using **PostgreSQL, Redis, Langfuse, and Temporal**, combining working, conversational, project, and user memory with token-aware compression and scoped retrieval to reduce P95 context-assembly latency from **1.2s to 430ms**
- Extended the Harness into a self-improvement loop, compiling user corrections and Langfuse-traced failures into versioned Skill candidates gated by sandbox replay, canary release, and rollback, cutting regressions on previously solved tasks from 21% to 5%

NIO

Shanghai, China

Machine learning Engineer Intern

May 2024 – Jul. 2024

- Built a pairwise preference dataset and performed **supervised fine-tuning (SFT) using LoRA (Low-Rank Adaptation, r=16)** followed by **DPO (Direct Preference Optimization)** in **PyTorch**, with senior inspectors, improving the defect detection rate from **76% to 91%**
- Rebuilt the automated evaluation pipeline with a continuously generated golden set, reducing human adjudication by moving to an **OpenCV/YOLO** model, improving annotation agreement from **71% to 94%**, and reducing labeling effort by **55%**

Capgemini

Shanghai, China

Machine Learning Engineer Intern

Aug 2024 – Sep. 2024

- Created a **Java/Spring Boot** transaction pipeline for atomic reservations, handling **1,200 QPS** with zero error tolerance
- Designed a **Redis/RocketMQ/MySQL** consistency layer for reliable delivery and idempotent processing, recovering from **95% of database failures** and reducing recovery time from **30 to 3 minutes**
- Engineered a **Redis** layer with hot-key protection, reducing peak MySQL reads by **75%** and improving P99 latency by **60%**
- Developed a reusable **Java** execution engine with **CompletableFuture** and isolated thread pools for parallel retrieval, reducing multi-source aggregation latency by **50%** and repetitive asynchronous code by **65%**

SKILLS

- **Languages:** Python | Java | SQL | TypeScript | JavaScript | Bash
- **AI & LLM:** LLM application development | RAG | Multi-agent systems | AI agents | OpenAI API | LangChain | LangGraph | Hugging Face | Fine-tuning (LoRA, SFT, DPO) | Prompt engineering | Embeddings & vector search | Text-to-SQL | LLM evaluation
- **Machine Learning:** PyTorch | scikit-learn | NumPy | Pandas | Computer vision | YOLO | OpenCV | TensorRT | Model training & deployment
- **Backend:** FastAPI | Spring Boot | REST APIs | Microservices | Distributed systems | High-concurrency systems | Asynchronous programming
- **Data & Messaging:** PostgreSQL | MySQL | Redis | Elasticsearch | pgvector | RocketMQ | Message queues | Object storage (MinIO, S3) | Data pipelines
- **Cloud & DevOps:** Docker | Kubernetes | Azure | Linux | Git | GitHub Actions | CI/CD | Monitoring & logging